

## Task A3.1

### Commands to create metrics-server and verify it works

1. Run “`kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml`” to download the metrics-server from online and apply it to our clusters.

```
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml
serviceaccount/metrics-server created
clusterrole.rbac.authorization.k8s.io/system:aggregated-metrics-reader created
clusterrole.rbac.authorization.k8s.io/system:metrics-server created
rolebinding.rbac.authorization.k8s.io/metrics-server-auth-reader created
clusterrolebinding.rbac.authorization.k8s.io/metrics-server:system:auth-delegator created
clusterrolebinding.rbac.authorization.k8s.io/system:metrics-server created
service/metrics-server created
deployment.apps/metrics-server created
apiservice.apiregistration.k8s.io/v1beta1.metrics.k8s.io created
```

2. Execute “`kubectl -nkube-system edit deploy/metrics-server`” to edit the metrics-server and add the “`--kubelet-insecure-tls`” flag to `deployment.spec.containers[].args[]`. Below we can see how I added it.

```
spec:
  containers:
  - args:
    - --cert-dir=/tmp
    - --secure-port=4443
    - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
    - --kubelet-use-node-status-port
    - --metric-resolution=15s
    - --kubelet-insecure-tls
    image: k8s.gcr.io/metrics-server/metrics-server:v0.6.1
    imagePullPolicy: IfNotPresent
    livenessProbe:
      failureThreshold: 3
      httpGet:
        path: /livez
        port: https
        scheme: HTTPS
      periodSeconds: 10
      successThreshold: 1
```

3. Run “`kubectl -nkube-system rollout restart deploy/metrics-server`” to restart metrics-server.

```
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl -nkube-system rollout restart deploy/metrics-server
deployment.apps/metrics-server restarted
```

### Commands to create HPA and verify it works

1. Run “`kubectl apply -f k8s/manifests/k8/backend-hpa.yaml`” to apply the HPA.

```
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl apply -f k8s/manifests/k8/backend-hpa.yaml
horizontalpodautoscaler.autoscaling/backend created
```

2. To test the HPA to see if it works, I tried stressing the system in different ways
  - a. I initially only queried the app intermittently to not stress the servers too much. As seen in screenshot below, as the load is below target, and the system does not need to scale up (see number of deployed pods and desired pods – stays constant).

```
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl describe hpa
Warning: autoscaling/v2beta2 HorizontalPodAutoscaler is deprecated in v1.23+, unavailable in v1.26+; use autoscaling/v2 HorizontalPodAutoscaler
Name: backend
Namespace: default
Labels: <none>
Annotations: <none>
CreationTimestamp: Sun, 30 Oct 2022 12:10:20 +0800
Reference: Deployment/backend
Metrics:
  resource cpu on pods (as a percentage of request): 15% (6m) / 50%
Min replicas: 1
Max replicas: 10
Deployment pods: 1 current / 1 desired
Conditions:
  Type            Status  Reason                        Message
  ----            -
  AbleToScale     True    ReadyForNewScale             recommended size matches current size
  ScalingActive   True    ValidMetricFound             the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
  ScalingLimited  False   DesiredWithinRange           the desired count is within the acceptable range
Events:
  Type    Reason              Age   From                      Message
  ----    -
  Normal  SuccessfulRescale   3m56s horizontal-pod-autoscaler  New size: 1; reason: All metrics below target
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl describe hpa
```

- b. I then queried very the app frequently to stress the servers above the target. As seen in the screenshot below, as the load increased above the target, the system scaled up (see number of deployed pods and desired pods – wants more pods).

```
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl describe hpa
Warning: autoscaling/v2beta2 HorizontalPodAutoscaler is deprecated in v1.23+, unavailable in v1.26+; use autoscaling/v2 HorizontalPodAutoscaler
Name: backend
Namespace: default
Labels: <none>
Annotations: <none>
CreationTimestamp: Sun, 30 Oct 2022 12:10:20 +0800
Reference: Deployment/backend
Metrics:
  resource cpu on pods (as a percentage of request): 80% (32m) / 50%
Min replicas: 1
Max replicas: 10
Deployment pods: 1 current / 2 desired
Conditions:
  Type            Status  Reason                        Message
  ----            -
  AbleToScale     True    SucceededRescale             the HPA controller was able to update the target scale to 2
  ScalingActive   True    ValidMetricFound             the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
  ScalingLimited  False   DesiredWithinRange           the desired count is within the acceptable range
Events:
  Type    Reason              Age   From                      Message
  ----    -
  Normal  SuccessfulRescale   4m20s horizontal-pod-autoscaler  New size: 1; reason: All metrics below target
  Normal  SuccessfulRescale   5s    horizontal-pod-autoscaler  New size: 2; reason: cpu resource utilization (percentage of request) above target
```

- c. I then stopped querying the app frequently to destress the servers below the target. As seen in the screenshot below, as the load is decreased below the target, the system scaled down (see number of deployed pods and desired pods – wants less).

```
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl describe hpa
Warning: autoscaling/v2beta2 HorizontalPodAutoscaler is deprecated in v1.23+, unavailable in v1.26+; use autoscaling/v2 HorizontalPodAutoscaler
Name: backend
Namespace: default
Labels: <none>
Annotations: <none>
CreationTimestamp: Sun, 30 Oct 2022 12:10:20 +0800
Reference: Deployment/backend
Metrics:
  resource cpu on pods (as a percentage of request): 3% (1m) / 50%
Min replicas: 1
Max replicas: 10
Deployment pods: 5 current / 4 desired
Conditions:
  Type            Status  Reason                        Message
  ----            -
  AbleToScale     True    SucceededRescale             the HPA controller was able to update the target scale to 4
  ScalingActive   True    ValidMetricFound             the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
  ScalingLimited  False   DesiredWithinRange           the desired count is within the acceptable range
Events:
  Type    Reason              Age   From                      Message
  ----    -
  Normal  SuccessfulRescale   14m   horizontal-pod-autoscaler  New size: 1; reason: All metrics below target
  Normal  SuccessfulRescale   9m46s horizontal-pod-autoscaler  New size: 2; reason: cpu resource utilization (percentage of request) above target
  Normal  SuccessfulRescale   6m30s horizontal-pod-autoscaler  New size: 3; reason: cpu resource utilization (percentage of request) above target
  Normal  SuccessfulRescale   6m    horizontal-pod-autoscaler  New size: 5; reason: cpu resource utilization (percentage of request) above target
  Normal  SuccessfulRescale   0s    horizontal-pod-autoscaler  New size: 4; reason: All metrics below target
```

3. We can confirm that the pods are running in different zones by executing “kubectl get nodes -L topology.kubernetes.io/zone”

```
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl get nodes -L topology.kubernetes.io/zone
```

NAME	STATUS	ROLES	AGE	VERSION	ZONE
kind-1-control-plane	Ready	control-plane	20h	v1.25.3	
kind-1-worker	Ready	<none>	20h	v1.25.3	a
kind-1-worker2	Ready	<none>	20h	v1.25.3	a
kind-1-worker3	Ready	<none>	20h	v1.25.3	b

## Task A3.2

### Commands to create the Deployment and verify it works

1. We apply the zone aware deployment with the command “kubectl apply -f k8s/manifests/k8/backend-deployment-zone-aware.yaml”

```
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl apply -f k8s/manifests/k8/backend-deployment-zone-aware.yaml
deployment.apps/backend-zone-aware created
```

2. We recall that “worker” and “worker2” are running in zone a while “worker3” is running in zone b. We run “kubectl get po -lapp=backend-zone-aware -o wide --sort-by='.spec.nodeName'” to check how many pods are assigned to each node. We can see that there are 4 pods assigned to worker, 1 to worker 2, and 5 to worker3. This means that there are 5 pods running in zone a and 5 pods in zone b, and thus we know that the pods are uniformly distributed across the two zones.

```
roberto@matebook:~/NUS/SEPP/OTOT/A2-A3$ kubectl get po -lapp=backend-zone-aware -o wide --sort-by='.spec.nodeName'
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
backend-zone-aware-5f879f6fb7-2n672	1/1	Running	0	94s	10.244.1.6	kind-1-worker	<none>	<none>
backend-zone-aware-5f879f6fb7-b4ktj	1/1	Running	0	94s	10.244.1.8	kind-1-worker	<none>	<none>
backend-zone-aware-5f879f6fb7-mdkj4	1/1	Running	0	94s	10.244.1.7	kind-1-worker	<none>	<none>
backend-zone-aware-5f879f6fb7-qdlch	1/1	Running	0	94s	10.244.1.9	kind-1-worker	<none>	<none>
backend-zone-aware-5f879f6fb7-2zlwf	1/1	Running	0	94s	10.244.3.6	kind-1-worker2	<none>	<none>
backend-zone-aware-5f879f6fb7-g4q8z	1/1	Running	0	94s	10.244.2.9	kind-1-worker3	<none>	<none>
backend-zone-aware-5f879f6fb7-hb6qs	1/1	Running	0	94s	10.244.2.10	kind-1-worker3	<none>	<none>
backend-zone-aware-5f879f6fb7-kw7dj	1/1	Running	0	94s	10.244.2.7	kind-1-worker3	<none>	<none>
backend-zone-aware-5f879f6fb7-sfc9f	1/1	Running	0	94s	10.244.2.11	kind-1-worker3	<none>	<none>
backend-zone-aware-5f879f6fb7-xxgv2	1/1	Running	0	94s	10.244.2.8	kind-1-worker3	<none>	<none>